

Machine Intelligence

11. Classical Planning

Automated Sequential Decision Making on “Simple” Environments

Álvaro Torralba



AALBORG UNIVERSITET

Fall 2022

Thanks to Thomas D. Nielsen and Jörg Hoffmann for slide sources

Planning

Ambition:


Write one program (planner) that can solve all sequential decision-making problems.

How do we describe our problem to the planner?

- A *logical description* of the possible **states**
- A *logical description* of the **initial state** I
- A *logical description* of the **goal condition** G
- *logical description* of the set A of **actions** in terms of **preconditions** and **effects**

→ Solution (**plan**) = **sequence of actions** from A , transforming I into a state that satisfies G .

Example of a Planning Task



- **States:** Card positions (*position_Jspades=Qhearts*).
- **Actions:** Card moves (*move_Jspades_Qhearts_freecell4*).
- **Initial state:** Start configuration.
- **Goal states:** All cards “home”.
- **Solution:** Card moves solving this game.

Algorithmic Problems in Planning

Satisficing Planning

Input: A planning task Π .

Output: A plan for Π , or “unsolvable” if no plan for Π exists.

Optimal Planning

Input: A planning task Π .

Output: An *optimal* plan for Π , or “unsolvable” if no plan for Π exists.

→ The techniques successful for either one of these are almost disjoint. And satisficing planning is much more effective in practice.

→ Programs solving these problems are called (optimal) **planners**, **planning systems**, or **planning tools**.

Our Agenda for This Chapter

- **The STRIPS Planning Formalism:** How do we represent a planning task?
→ **Lays the framework we'll be looking at.**
- **Planning Domain Definition Language:** How we actually express our classical planning problems.
→ **For reference.**
- **Planning as Heuristic Search:** How state-space search techniques are applied to solve planning tasks?
→ **A Recap of Chapter 2.**
- **The Delete Relaxation:** How to relax a planning problem?
→ **The delete relaxation is the most successful method for the automatic generation of heuristic functions. It is a key ingredient of many IPC winners during the last two decades. It relaxes STRIPS planning tasks by ignoring the delete lists.**
- **The h^+ Heuristic:** What is the resulting heuristic function?
→ **h^+ is the “ideal” delete relaxation heuristic.**
- **Approximating h^+ :** How to actually compute a heuristic?
→ **Turns out that, in practice, we must approximate h^+ .**

“STRIPS” Planning

- **STRIPS** = Stanford Research Institute Problem Solver.

STRIPS is the simplest possible (reasonably expressive) logics-based planning language.

- STRIPS has only **Boolean variables**: propositional logic atoms.
- Its preconditions/effects/goals are as canonical as imaginable:
 - Preconditions, goals: **conjunctions** of **positive atoms**.
 - Effects: **conjunctions** of **literals** (positive or negated atoms).
- We use the common set-based notation for this simple formalism.

→ Historical note: STRIPS [?] was originally a planner, whose language actually wasn't quite that simple.

STRIPS Planning: Syntax

Definition (STRIPS Planning Task). A **STRIPS planning task**, short **planning task**, is a 4-tuple $\Pi = (P, A, I, G)$ where:

- P is a finite set of **facts** (aka **propositions**).
- A is a finite set of **actions**; each $a \in A$ is a triple $a = (pre_a, add_a, del_a)$ of subsets of P referred to as the action's **precondition**, **add list**, and **delete list** respectively; we require that $add_a \cap del_a = \emptyset$.
- $I \subseteq P$ is the **initial state**.
- $G \subseteq P$ is the **goal**.

We will often give each action $a \in A$ a **name** (a string), and identify a with that name.

Note: We assume **unit costs** for simplicity: every action has cost 1.

“TSP” in Australia



STRIPS Encoding of “TSP”



- **Facts** P : $\{at(x), visited(x) \mid x \in \{Sydney, Adelaide, Brisbane, Perth, Darwin\}\}$.
- **Initial state** I :
- **Goal** G :
- **Actions** $a \in A$: $drive(x, y)$ where x, y have a road.
Precondition pre_a :
Add list add_a :
Delete list del_a :
- **Plan**:

STRIPS Planning: Semantics

Definition (STRIPS State Space). Let $\Pi = (P, A, c, I, G)$ be a STRIPS planning task. The **state space** of Π is $\Theta_\Pi = (S, A, T, I, S^G)$ where:

- The states (also **world states**) $S = 2^P$ are the subsets of P .
- A is Π 's action set.
- The transitions are $T = \{s \xrightarrow{a} s' \mid \text{pre}_a \subseteq s, s' = \text{appl}(s, a)\}$.
If $\text{pre}_a \subseteq s$, then a is **applicable** in s and $\text{appl}(s, a) := (s \cup \text{add}_a) \setminus \text{del}_a$. If $\text{pre}_a \not\subseteq s$, then $\text{appl}(s, a)$ is undefined.
- I is Π 's initial state.
- The goal states $S^G = \{s \in S \mid G \subseteq s\}$ are those that satisfy Π 's goal.

An (optimal) **plan** for $s \in S$ is an (optimal) solution for s in Θ_Π , i.e., a path from s to some $s' \in S^G$. A solution for I is called a **plan for Π** . Π is **solvable** if a plan for Π exists.

For $\vec{a} = \langle a_1, \dots, a_n \rangle$, $\text{appl}(s, \vec{a}) := \text{appl}(\dots \text{appl}(\text{appl}(s, a_1), a_2) \dots, a_n)$ if each a_i is applicable in the respective state; else, $\text{appl}(s, \vec{a})$ is undefined.

STRIPS Encoding of Simplified “TSP”




- **Facts** P : $\{at(x), visited(x) \mid x \in \{Sydney, Adelaide, Brisbane\}\}$.
- **Initial state** I :
- **Goal** G : $\{visited(x) \mid x \in \{Sydney, Adelaide, Brisbane\}\}$. (Note: no “ $at(Sydney)$ ”.)
- **Actions** $a \in A$: $drive(x, y)$ where x, y have a road.
 - Precondition** pre_a :
 - Add list** add_a :
 - Delete list** del_a :

STRIPS Encoding of Simplified “TSP”: State Space

→ Is this actually the state space?

(Oh no it's) The Blocksworld



- **Facts:** $on(x, y)$, $onTable(x)$, $clear(x)$, $holding(x)$, $armEmpty()$.
- **Initial state:** $\{onTable(E), clear(E), \dots, onTable(C), on(D, C), clear(D), armEmpty()\}$.
- **Goal:** $\{on(E, C), on(C, A), on(B, D)\}$.
- **Actions:** $stack(x, y)$, $unstack(x, y)$, $putdown(x)$, $pickup(x)$.
- $stack(x, y)?$

Questionnaire

Question!

Which are correct encodings (part of some correct overall encoding) of the STRIPS Blocksworld *pickup*(*x*) action schema?

(A): ($\{onTable(x), clear(x), armEmpty()\}, \{holding(x)\}, \{onTable(x)\}$).

(C): ($\{onTable(x), clear(x), armEmpty()\}, \{holding(x)\}, \{onTable(x), armEmpty(), clear(x)\}$).

(B): ($\{onTable(x), clear(x), armEmpty()\}, \{holding(x)\}, \{armEmpty()\}$).

(D): ($\{onTable(x), clear(x), armEmpty()\}, \{holding(x)\}, \{onTable(x), armEmpty()\}$).


Levels of Representation (Chapter 1)

We consider 3 representation schemes

State based:

State-space Search

(Chapter 2)




State 01

Feature based:

STRIPS


(This Chapter)




*egg=whole,
butter_in=table,
egg_in=table*

Relational: PDDL


(What we actually use,
not in this course)



*state(egg,whole),
in(butter,table),
in(egg,table)*



State 12




*egg=broken,
butter_in=table,
egg_in=pan*




*state(egg,broken),
in(butter,table),
in(egg,pan)*

The Blocksworld in PDDL (STRIPS):



Initial State



Goal State

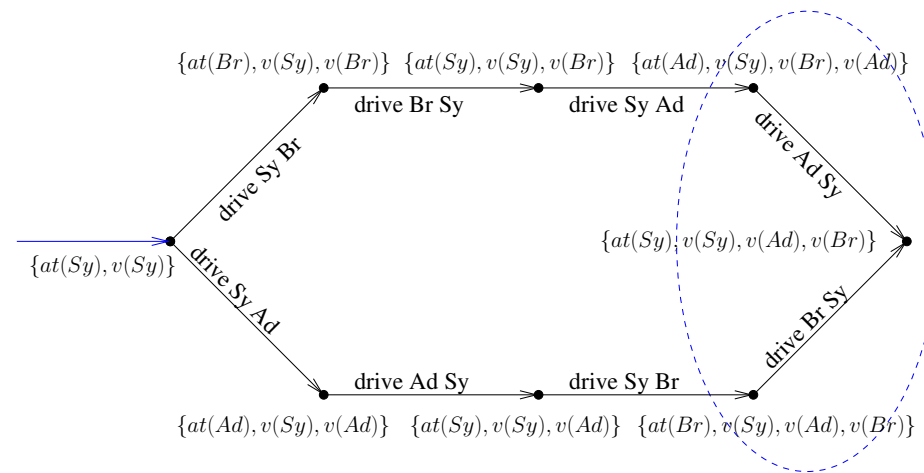
Domain File:

```
(define (domain blocksworld)
  (:predicates (clear ?x) (holding ?x) (on ?x ?y)
               (on-table ?x) (arm-empty))
  (:action stack
    :parameters (?x ?y)
    :precondition (and (clear ?y) (holding ?x))
    :effect (and (arm-empty) (on ?x ?y)
                 (not (clear ?y)) (not (holding ?x))))
  )
  ...
```

Problem File:

```
(define (problem bw-abcde)
  (:domain blocksworld)
  (:objects a b c d e)
  (:init (on-table a) (clear a)
         (on-table b) (clear b)
         (on-table e) (clear e)
         (on-table c) (on d c) (clear d)
         (arm-empty))
  (:goal (and (on e c) (on c a) (on b d))))
```


So, How to Solve Planning Problems?




Question

Given a graph G , a node I , and a set of nodes G , find the shortest path from I to any node in G . What algorithm do you suggest to use?

The Problem


Example: Planning as Search

Example: “Logistics”




- **Facts P :** $\{truck(x) \mid x \in \{A, B, C, D\}\} \cup \{pack(x) \mid x \in \{A, B, C, D, T\}\}$.
- **Initial state I :** $\{truck(A), pack(C)\}$.
- **Goal G :** $\{truck(A), pack(D)\}$.
- **Actions A :** (Notated as “precondition \Rightarrow adds, \neg deletes”)
 - $drive(x, y)$, where x, y have a road: “ $truck(x) \Rightarrow truck(y), \neg truck(x)$ ”.
 - $load(x)$: “ $truck(x), pack(x) \Rightarrow pack(T), \neg pack(x)$ ”.
 - $unload(x)$: “ $truck(x), pack(T) \Rightarrow pack(x), \neg pack(T)$ ”.

Example: Planning as Search



Heuristic Search



→ Heuristic function h estimates the cost of an optimal path from a state s to the goal; search prefers to expand states s with small $h(s)$.

Definition (Heuristic Function). Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$. A **heuristic function**, short **heuristic**, for Π is a function $h : S \mapsto \mathbb{R}_0^+ \cup \{\infty\}$. Its value $h(s)$ for a state s is referred to as the state's **heuristic value**, or **h value**.

Definition (Remaining Cost, h^*). Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$. For a state $s \in S$, the state's **remaining cost** is the cost of an optimal plan for s , or ∞ if there exists no plan for s . The **perfect heuristic** for Π , written h^* , assigns every $s \in S$ its remaining cost as the heuristic value.

→ Heuristic functions h estimate remaining cost h^* .

Properties of Individual Heuristic Functions

Definition (Safe/Goal-Aware/Admissible/Consistent). Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$, and let h be a heuristic for Π . The heuristic is called:

- **safe** if, for all $s \in S$, $h(s) = \infty$ implies $h^*(s) = \infty$;
- **goal-aware** if $h(s) = 0$ for all goal states $s \in S^G$;
- **admissible** if $h(s) \leq h^*(s)$ for all $s \in S$;
- **consistent** if $h(s) \leq h(s') + c(a)$ for all transitions $s \xrightarrow{a} s'$.

→ Relationships:

Proposition. Let Π be a planning task, and let h be a heuristic for Π . Then:

- If h is admissible, then h is goal-aware.
- If h is admissible, then h is safe.
- If h is consistent and goal-aware, then h is admissible.
- No other implications of this form hold.

Greedy Best-First Search and A^*

Duplicate elimination omitted

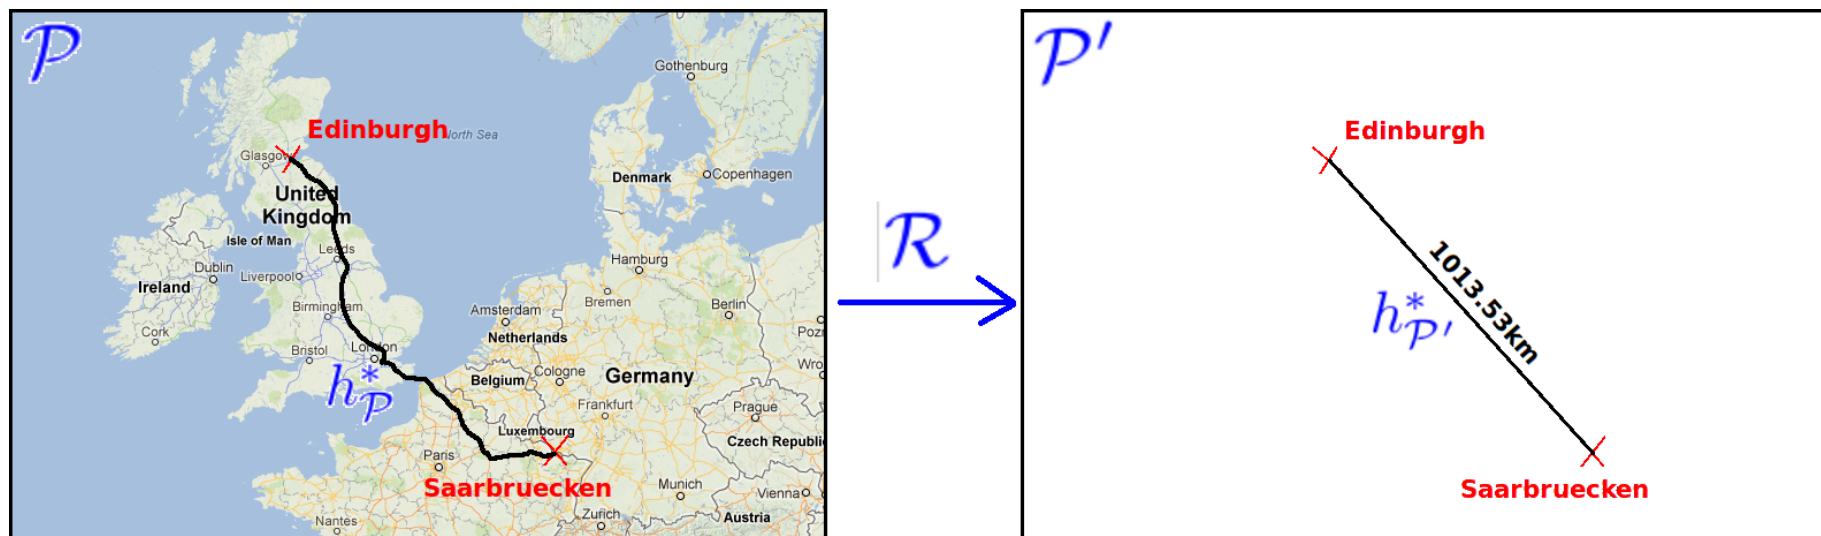
```
function Greedy Best-First Search [ $A^*$ ](problem) returns a solution, or failure
  node  $\leftarrow$  a node  $n$  with  $n.state = problem.InitialState$ 
  frontier  $\leftarrow$  a priority queue ordered by ascending  $h$  [ $g + h$ ], only element  $n$ 
  loop do
    if Empty?(frontier) then return failure
     $n \leftarrow Pop(frontier)$ 
    if problem.GoalTest( $n.State$ ) then return Solution( $n$ )
    for each action  $a$  in problem.Actions( $n.State$ ) do
       $n' \leftarrow ChildNode(problem, n, a)$ 
      Insert( $n'$ ,  $h(n')$  [ $g(n') + h(n')$ ], frontier)
```

→ Greedy best-first search explores states by increasing heuristic value h . A^* explores states by increasing plan-cost estimate $g + h$.

Greedy best-first search: **Fast but not optimal \implies satisficing planning.**

A^* : **Optimal for admissible $h \implies$ optimal planning,
with such h .**


Relaxation in Route-Finding



- Problem class \mathcal{P} : Route finding.
- Perfect heuristic $h_{\mathcal{P}}^*$ for \mathcal{P} : Length of a shortest route.
- Simpler problem class \mathcal{P}' :
- Perfect heuristic $h_{\mathcal{P}'}^*$ for \mathcal{P}' :
- Transformation \mathcal{R} :

How to Relax in Planning? (A Reminder!)

Example: “Logistics”



- **Facts P :** $\{truck(x) \mid x \in \{A, B, C, D\}\} \cup \{pack(x) \mid x \in \{A, B, C, D, T\}\}$.
- **Initial state I :** $\{truck(A), pack(C)\}$.
- **Goal G :** $\{truck(A), pack(D)\}$.
- **Actions A :** (Notated as “precondition \Rightarrow adds, \neg deletes”)
 - $drive(x, y)$, where x, y have a road: “ $truck(x) \Rightarrow truck(y), \neg truck(x)$ ”.
 - $load(x)$: “ $truck(x), pack(x) \Rightarrow pack(T), \neg pack(x)$ ”.
 - $unload(x)$: “ $truck(x), pack(T) \Rightarrow pack(x), \neg pack(T)$ ”.

Example “Only-Adds” Relaxation: **Drop the preconditions and deletes.**

“ $drive(x, y): \Rightarrow truck(y)$ ”; “ $load(x): \Rightarrow pack(T)$ ”; “ $unload(x): \Rightarrow pack(x)$ ”.

→ **Heuristic value for I is?**

How to Relax During Search: Overview

Attention! Search uses the real (un-relaxed) Π . The relaxation is applied (e.g., in Only-Adds, the simplified actions are used) only within the call to $h(s)$!!!

- A common student mistake is to instead apply the relaxation once to the whole problem, then doing the whole search “within the relaxation”.
- The next slide illustrates the correct search process in detail.

How to Relax During Search: Only-Adds

How the Delete Relaxation Changes the World

The Delete Relaxation

Definition (Delete Relaxation). Let $\Pi = (P, A, I, G)$ be a planning task. The **delete-relaxation** of Π is the task $\Pi^+ = (P, A^+, I, G)$ where $A^+ = \{a^+ \mid a \in A\}$ with $pre_{a^+} = pre_a$, $add_{a^+} = add_a$, and $del_{a^+} =$

→ In other words, the class of simpler problems \mathcal{P}' is the set of all STRIPS planning tasks with empty delete lists, and the relaxation mapping \mathcal{R} drops the delete lists.

Definition (Relaxed Plan). Let $\Pi = (P, A, I, G)$ be a planning task, and let s be a state. A **relaxed plan** for s is a plan for s . A relaxed plan for I is called a relaxed plan for Π .

→ A relaxed plan for s is an action sequence that solves s when pretending that all delete lists are empty.


→ Also called **delete-relaxed plan**; “relaxation” is often used to mean “delete-relaxation” by default.

A Relaxed Plan for “TSP” in Australia



- 1 **Initial state:** $\{at(Sydney), visited(Sydney)\}$.
- 2 **Apply** $drive(Sydney, Brisbane)^+$: $\{at(Brisbane), visited(Brisbane), at(Sydney), visited(Sydney)\}$.
- 3 **Apply** $drive(Sydney, Adelaide)^+$: $\{at(Adelaide), visited(Adelaide), at(Brisbane), visited(Brisbane), at(Sydney), visited(Sydney)\}$.
- 4 **Apply** $drive(Adelaide, Perth)^+$: $\{at(Perth), visited(Perth), at(Adelaide), visited(Adelaide), at(Brisbane), visited(Brisbane), at(Sydney), visited(Sydney)\}$.
- 5 **Apply** $drive(Adelaide, Darwin)^+$: $\{at(Darwin), visited(Darwin), at(Perth), visited(Perth), at(Adelaide), visited(Adelaide), at(Brisbane), visited(Brisbane), at(Sydney), visited(Sydney)\}$.

A Relaxed Plan for “Logistics”



- **Facts P :** $\{truck(x) \mid x \in \{A, B, C, D\}\} \cup \{pack(x) \mid x \in \{A, B, C, D, T\}\}$.
- **Initial state I :** $\{truck(A), pack(C)\}$.
- **Goal G :** $\{truck(A), pack(D)\}$.
- **Relaxed actions A^+ :** (Notated as “precondition \Rightarrow adds”)
 - $drive(x, y)^+$: “ $truck(x) \Rightarrow truck(y)$ ”.
 - $load(x)^+$: “ $truck(x), pack(x) \Rightarrow pack(T)$ ”.
 - $unload(x)^+$: “ $truck(x), pack(T) \Rightarrow pack(x)$ ”.

Relaxed plan:

Definition (Relaxed Plan Existence Problem). By **PlanEx⁺**, we denote the problem of deciding, given a planning task $\Pi = (P, A, I, G)$, whether or not there exists a **relaxed plan** for Π .

→ **This is easier than PlanEx for general STRIPS!**

Proposition (PlanEx⁺ is Easy). *PlanEx⁺ is a member of P.*


Proof. The following algorithm decides PlanEx⁺:

```

F := I
while G ⊄ F do
    F' := F ∪ ⋃a ∈ A: prea ⊆ F adda
    (*) if F' = F then return “unsolvable” endif
    F := F'
endwhile
return “solvable”
    
```

The algorithm terminates after at most


Deciding PlanEx⁺ in “Logistics”



Iterations on F :

- $\{truck(A), pack(C)\}$
- \cup
- \cup
- \cup
- \cup

Deciding PlanEx⁺ in Unsolvable “Logistics”



Iterations on F :

- $\{truck(A), pack(C)\}$
- \cup
- \cup
- \cup
- \cup
- \cup

Questionnaire

Question!

How does ignoring delete lists simplify Sokoban?

(A): You will never “lock yourself in”.

(C): You can walk through walls.

(B): Free positions remain free.

(D): A single action can push 2 stones at once.

The h^+ Heuristic

→ PlanEx^+ is not actually what we're looking for. $\text{PlanEx}^+ = \text{relaxed plan existence}$; we want relaxed plan length $h^* \circ \mathcal{R}$.

Definition (Optimal Relaxed Plan). Let $\Pi = (P, A, I, G)$ be a planning task, and let s be a state. An **optimal relaxed plan** for s is an optimal plan for $(P, A, s, G)^+$.

Here's what we're looking for:

Definition (h^+). Let $\Pi = (P, A, I, G)$ be a planning task with states S . The **ideal delete-relaxation heuristic** h^+ for Π is the function $h^+ : S \mapsto \mathbb{N}_0 \cup \{\infty\}$ where $h^+(s)$ is the length of an optimal relaxed plan for s if a relaxed plan for s exists, and $h^+(s) = \infty$ otherwise.

→ In other words, $h^+ = h^* \circ \mathcal{R}$, cf. previous slide.

h^+ is Admissible

Lemma. Let $\Pi = (P, A, I, G)$ be a planning task, and let s be a state. If $\langle a_1, \dots, a_n \rangle$ is a plan for (P, A, s, G) , then $\langle a_1^+, \dots, a_n^+ \rangle$ is a plan for $(P, A, s, G)^+$.

Proof Sketch. (for reference) Show by induction over $0 \leq i \leq n$ that $appl(s, \langle a_1, \dots, a_i \rangle) \subseteq appl(s, \langle a_1^+, \dots, a_i^+ \rangle)$.

→ "If we ignore deletes, the states along the plan can only get bigger."

Theorem. h^+ is Admissible.

Proof. (for reference) Let $\Pi = (P, A, I, G)$ be a planning task with states S , and let $s \in S$. $h^+(s)$ is defined as optimal plan length in $(P, A, s, G)^+$. With the above lemma, any plan for (P, A, s, G) also constitutes a plan for $(P, A, s, G)^+$. Thus optimal plan length in $(P, A, s, G)^+$ cannot be longer than that in (P, A, s, G) , and the claim follows.

h^+ in “TSP” in Australia



Planning vs. Relaxed Planning:

- **Optimal plan:** $\langle \text{drive}(\text{Sydney}, \text{Brisbane}), \text{drive}(\text{Brisbane}, \text{Sydney}), \text{drive}(\text{Sydney}, \text{Adelaide}), \text{drive}(\text{Adelaide}, \text{Perth}), \text{drive}(\text{Perth}, \text{Adelaide}), \text{drive}(\text{Adelaide}, \text{Darwin}), \text{drive}(\text{Darwin}, \text{Adelaide}), \text{drive}(\text{Adelaide}, \text{Sydney}) \rangle$.
- **Optimal relaxed plan:** $\langle \text{drive}(\text{Sydney}, \text{Brisbane}), \text{drive}(\text{Sydney}, \text{Adelaide}), \text{drive}(\text{Adelaide}, \text{Perth}), \text{drive}(\text{Adelaide}, \text{Darwin}) \rangle$.
- $h^*(I) = 8; h^+(I) = 4$.

How to Relax During Search: Ignoring Deletes

Approximating h^+ : h^{FF}

Theorem. $PlanLen^+$ is **NP-complete**.

→ We can't compute our heuristic h^+ efficiently. So we approximate it instead.

Definition (h^{FF}). Let $\Pi = (P, A, I, G)$ be a planning task with states S . A **relaxed plan heuristic** h^{FF} for Π is a function $h^{FF} : S \mapsto \mathbb{N}_0 \cup \{\infty\}$ returning **the length of some, not necessarily optimal, relaxed plan for s** if a relaxed plan for s exists, and returning $h^{FF}(s) = \infty$ otherwise.

Notes:

- $h^{FF} \geq h^+$, i.e., h^{FF} never under-estimates h^+ .
- **We may have $h^{FF} > h^*$, i.e., h^{FF} is not admissible!** Thus h^{FF} can be used for satisficing planning only, not for optimal planning.

Observe: h^{FF} **as per this definition is not unique.** How do we find “some, not necessarily optimal, relaxed plan for (P, A, s, G) ”?

→ In what follows, we consider the following algorithm computing relaxed plans, and therewith (one variant of) h^{FF} :

- ① Chain **forward** to build a **relaxed planning graph (RPG)**.
- ② Chain **backward** to extract a relaxed plan from the RPG.

Computing h^{FF} : Relaxed Planning Graphs (RPG)

```
 $F_0 := s, t := 0$   
while  $G \not\subseteq F_t$  do  
   $A_t := \{a \in A \mid pre_a \subseteq F_t\}$   
   $F_{t+1} := F_t \cup \bigcup_{a \in A_t} add_a$   
  if  $F_{t+1} = F_t$  then stop endif  
   $t := t + 1$   
endwhile
```

→ Does this look familiar to you?

Computing h^{FF} : Extracting a Relaxed Plan

Information from the RPG: (min over an empty set is ∞)

- For $p \in P$: $level(\mathbf{p}) := \min\{\mathbf{t} \mid \mathbf{p} \in \mathbf{F}_t\}$.
- For $a \in A$: $level(\mathbf{a}) := \min\{\mathbf{t} \mid \mathbf{a} \in \mathbf{A}_t\}$.

```

M := max{level(p) | p ∈ G}
If M = ∞ then  $h^{\text{FF}}(s) := \infty$ ; stop endif
for t := 0, ..., M do
     $G_t := \{g \in G \mid level(g) = t\}$ 
endfor
for t := M, ..., 1 do
    for all g ∈  $G_t$  do
        select a, level(a) = t − 1, g ∈ adda
        for all p ∈ prea do  $G_{level(p)} := G_{level(p)} \cup \{p\}$  endfor
    endfor
endfor
 $h^{\text{FF}}(s) :=$  number of selected actions
  
```

Computing h^{FF} in “TSP” in Australia



RPG:

- $F_0 =$
- $A_0 =$
- $F_1 = F_0 \cup$
- $A_1 = A_0 \cup \{drive(Adelaide, Sydney), drive(Brisbane, Sydney)\}.$
- $F_2 = F_1 \cup$

Summary

- **Planning** is a form of general problem solving: develop solvers that perform well across a large class of problems.
- STRIPS is the simplest possible, while reasonably expressive, language for our purposes. It uses Boolean variables (facts), and defines actions in terms of precondition, add list, and delete list.
- We will consider Greedy Best-First Search and A^* as heuristic search algorithms.
- Heuristic search on classical search problems relies on a function h mapping states s to an estimate $h(s)$ of their goal distance. Such functions h are derived by solving **relaxed problems**.
- In planning, the relaxed problems are generated and solved automatically.
- The **delete relaxation** consists in dropping the deletes from STRIPS planning tasks. A **relaxed plan** is a plan for such a relaxed task. $h^+(s)$ is the length of an optimal relaxed plan for state s .

On the “Accuracy” of h^+

Reminder: Heuristics based on ignoring deletes are the key ingredient to almost all winners of the International Planning Competition in the last two decades.

→ **Why?**

→ A heuristic function is useful if its estimates are “accurate”.

How to measure this?

- **Known method 1:** Error relative to h^* , i.e., bounds on $|h^*(s) - h(s)|$.
- **Known method 2:** Properties of the **search space surface**: Local minima etc.

→ **For h^+ , method 2 is the road to success:**


→ In many benchmarks, under h^+ , local minima *provably* do not exist! [?]

A Brief Glimpse of h^+ Search Space Surfaces

h^+ in (the Real) TSP

h^+ in Graphs

Questionnaire



Question!

In this domain, h^+ is equal to?

(A): Manhattan Distance.

(B): Horizontal distance.

(C): Vertical distance.

(D): h^* .