

# Machine Intelligence

## 13. Planning under Uncertainty: Markov Decision Processes

Álvaro Torralba



**AALBORG UNIVERSITET**

Fall 2022

Thanks to Thomas D. Nielsen and Jörg Hoffmann for slide sources

## In previous chapters of the MI Lecture...

- State-space search **Chapter 2**:  
→ The basic algorithms under the hood
- Classical Planning **Chapter 11**:  
→ Automated Decision Making with domain-independent techniques. The user only specifies how the environment works and by analyzing the environment (delete-relaxation in our case), we can solve it efficiently!


Those techniques give us a **plan to achieve a goal** on environments that are:

- single-agent → we studied multi-agent environments on **Chapter 12**
- deterministic → **(This Chapter)**: non-determinism (MDPs)
- goal-oriented → **(This Chapter)**: reward-oriented (MDPs)
- fully-observable → partially-observable (POMDPs) (not in this lecture)
- discrete
- ...

# Decision problems with an unbounded time horizon

## Characteristics.

- at each step we are faced with a decision,
- at each step we are given a certain reward (possibly negative) determined by the chosen decision and the state of the world,
- the outcome of a decision may be uncertain (we will assume we know the probabilities associated to each outcome),
- the time horizon of the decision problem is unbounded.



# Markov Decision Processes

**Definition (Markov Decision Process).** An **MDP** is a 6-tuple  $\Theta = (S, A, r, P, I, S^T)$  where:




- $S$  is a finite set of **states**.
- $A$  is a finite set of **actions**.
- $r : S \times A \times S \mapsto \mathbb{R}$  is the **reward function**.
- $P : S \times A \times S \mapsto [0, 1]$  is the **transition probability function**, representing  $P(s' \mid s, a)$ , the probability that by applying action  $a$  in state  $s$ , we end in state  $s'$ .
- $I \in S$  is the **initial state**.
- $S^T \subseteq S$  is the set of **terminal states**. The set of terminal states could be empty.

→ We select actions in order to maximize the reward obtained

Similar to the definition of state-space (**c.f. Chapters 2/11**), but:

- Instead of a cost function, we have a reward function.
- Instead of a transition relation, we have a probability distribution over possible outgoing states

# Markov Decision Process Example

	1	2	3
1			
2			
3			




- The adventurer wants to take the gold located in (3,1).
- But she should avoid the snake.
- She has a complete knowledge of the environment.
- But paths are slippery and there is a small probability of failing one step and moving to the cell to the right.

To determine what's the best course of action, let's encode the problem as an MDP!

# Transition Probabilities in an MDP

The **transition probabilities**  $P(S_i | A_{i-1}, S_{i-1})$  tell us what are the possible outcomes of executing each action on each state

- The adventurer can move north, east, south, and west.
- A move succeeds with probability 0.8; otherwise it moves to the location immediately to the right.
- An attempt to move outside the boundaries of the grid results in staying in the current square.



	1	2	3
1			
2			
3			

**What are the outcomes of moving north in the initial state?**

For example, for the north action we have that  $P(S_{i+1} | \text{north}, S_i)$ :

(column,row)		$S_i$								
		(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	(3,1)	(3,2)	(3,3)
$S_{i+1}$	(1,1)	0.8	0.8	0	0	0	0	0	0	0
	(1,2)	0	0	0.8	0	0	0	0	0	0
	(1,3)	0	0	0	0	0	0	0	0	0
	(2,1)	0.2	0	0	0.8	0.8	0	0	0	0
	(2,2)	0	0.2	0	0	0	0.8	0	0	0
	(2,3)	0	0	0.2	0	0	0	0	0	0
	(3,1)	0	0	0	0.2	0	0	1	0.8	0
	(3,2)	0	0	0	0	0.2	0	0	0	0.8
	(3,3)	0	0	0	0	0	0.2	0	0	0

# Markov Decision Process as a Graph




- The state space of an MDP is an hypergraph where each action may have more than one possible outcome
- Similar to MinMax trees from Games, but here we have the probability of each outcome, instead of always picking the worst-case scenario
- The picture in this slide only shows a small portion (the actions applicable in the initial state): to generate the full state-space of the MDP we need to consider all actions in all states.

# Rewards

The **reward function**  $r(s, a, s')$  specifies “how many points we get” every time we apply an action.

In our example:

- Getting gold: +10
- Meeting snake: -5
- Moving elsewhere: -0.1 (so that shorter paths are preferred)



1	2	3

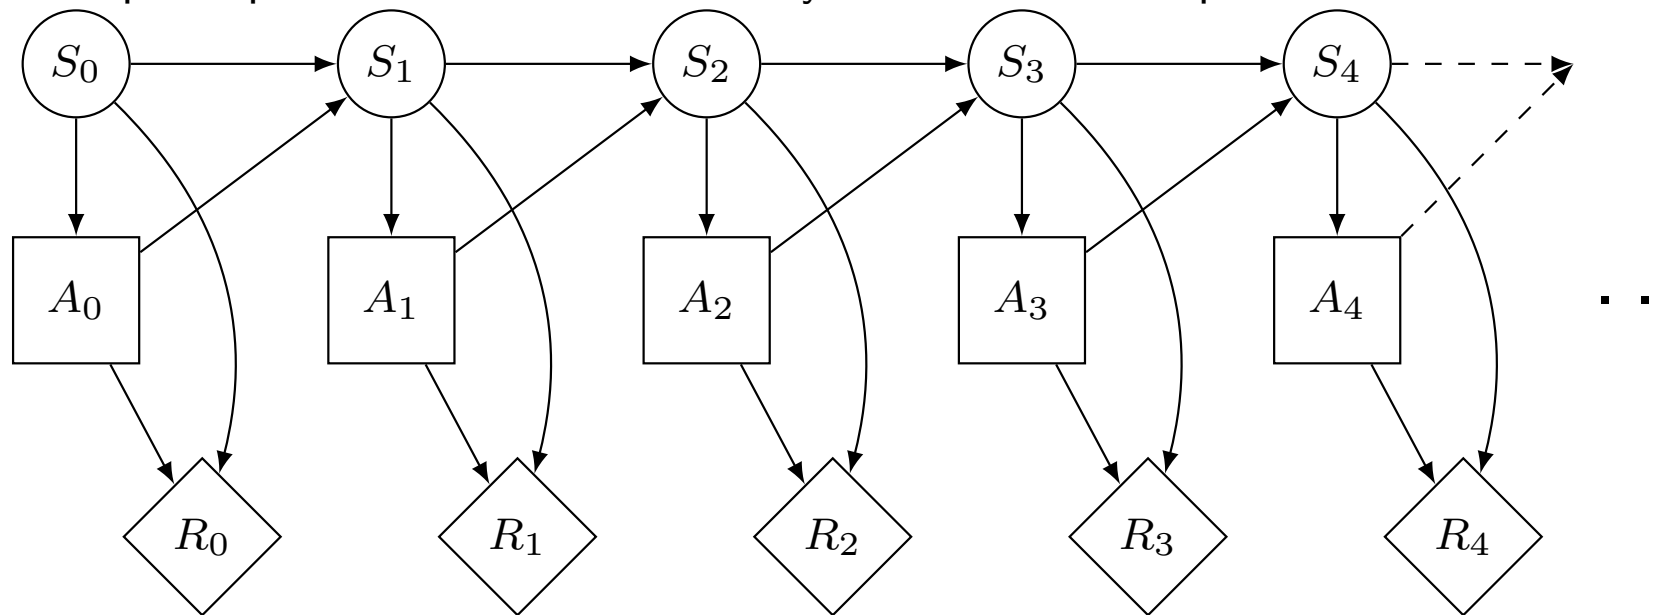


# Markov Assumption




**Markov Assumption:** given the present, the future does not depend on the past  
Where do we see the Markov assumption in Markov Decision Processes?

- Transition probabilities  $P(S_i | A_{i-1}, S_{i-1})$   
→ The probability that we reach certain state only depends on the present (what state we were and what action we applied)
- Reward function  $r(s, a, s')$   
→ Again, it does only depend on the present

Example of problem that does not satisfy the Markov assumption?



# Solving MDPs

	1	2	3
1			
2			
3			

## Question

What's the plan?

## Decision policies

A decision policy for MDPs is a function mapping non-terminal states to actions:

$$\pi : S \mapsto A$$

	1	2	3
1	→	→	
2	↑	↑	↑
3	↑	→	↑

In MDPs it suffices to consider this “simple” pure policies. In other settings policies could be more complex:

- A policy could depend on the entire history from the initial state to  $s$   
 $(\pi(s_0, a_0, s_1, a_1, \dots, s_i) \mapsto A)$   
 → We do not need to take the past into account due to the Markov property  
 In other words, if we visit the same state twice, we'd like to do the same action
- A policy could be mixed (assigning probabilities to actions)  
 → We do not need that here (in MDPs there is always some optimal policy that is pure).

## Executing a Policy


Given a policy, we can simulate its execution to obtain all possible traces it could have and their corresponding probability and utility.

Two example traces:

	1	2	3
1	→	→	
2	↑	↑	↑
3	↑	→	↑

# Markov Chains

The **Markov chain** associated to a policy  $\pi$  of an MDP is the (possibly infinite) set of traces with their corresponding probabilities.



The **expected utility** is the average reward obtained from each trace (weighted by its probability)

# Evaluating strategies

Imagine that we have the following reward function and **there is no terminal state** and no uncertainty on the result of an action:

	1	2	3
1	-0.1	-0.1	10
2	-0.1	-0.1	-5
3	-0.1	-0.1	-0.1

## Question!

Which strategy is best according to expected utility with unbounded horizon?

(A):  $\pi_A =$

	1	2	3
1	→	→	←
2	↑	↑	↑
3	↑	→	↑

(B):  $\pi_B =$

	1	2	3
1	→	→	↓
2	↑	↑	↓
3	↑	←	←

## Beyond Expected Sum of Rewards

**Problem:** Sum of rewards over an infinite horizon is often infinite, no matter how much you wait to start gaining reward

Two solutions:

- **Finite Horizon**
- **Discounted Rewards**

# The utility of an unbounded sequence: fixed horizon

## Finite-Horizon MDP

Consider only the rewards obtained in the first  $k$  states:

$$U(s_0, s_1, s_2, \dots) = R(s_0) + R(s_1) + \dots + R(s_k) < \infty.$$

But how do we choose  $k$ ?

- For  $k = 0$  we only care about the immediate reward, hence we pursue a very greedy strategy.
- For, say,  $k = 5$ , we only care about maximizing the reward the next 5 time steps (but we may end in a really bad state for the future)

**Finite horizon MDPs = Live your life as if there were not tomorrow**

Question!

Do finite-horizon MDPs satisfy the Markov property?

(A): Yes

(B): No



# Discounted rewards

Compare reward sequences

−0.1	10.0	−0.1	10.0	−0.1	10.0	−0.1	10.0	−0.1	10.0	−0.1	...
−0.1	−0.1	−0.1	−0.1	10.0	−0.1	−0.1	−0.1	−0.1	10.0	−0.1	...

## The utility of an unbounded sequence: discounted rewards

Weigh rewards in the immediate future higher than rewards in the distant future:

$$U(s_0, s_1, s_2, \dots) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots,$$

where  $0 \leq \gamma \leq 1$ . Possible interpretations of the **discounting factor**  $\gamma$ :

- In economics,  $\gamma$  may be thought of as inflation or an interest rate.
- The decision process may terminate with probability  $(1 - \gamma)$  at any point in time, e.g. the robot breaking down.

Thanks to increasing the exponent in gamma after every step, the sum of rewards is always finite!

With  $\gamma < 1$  we have:

$$U(s_0, s_1, s_2, \dots) = \sum_{i=0}^{\infty} \gamma^i R(s_i) \leq \sum_{i=0}^{\infty} \gamma^i \max R = \frac{\max R}{1 - \gamma} < \infty.$$

- For  $\gamma = 0$  we have a greedy strategy.
- For  $\gamma = 1$  we have normal additive rewards.

→ So, **expected discounted reward** is a good metric to compare policies → Typically, we want values  $\gamma < 1$  but  $\gamma \approx 1$ , such as 0.99

## Expected utilities

The actions may be non-deterministic so a strategy may only take you to a state with a certain probability.



Strategies should be compared based on the expected rewards they can produce.

Starting in state  $s_0$  and following strategy  $\pi$ , the **expected (discounted) reward in step  $i$**  is:

$$\underbrace{\gamma^i}_{\text{discount}} \sum_{s_i} R(s_i) \underbrace{P(S_i = s_i \mid \pi, S_0 = s_0)}_{\text{prob}(*)}$$

$\text{prob} (*)$  is the probability of reaching  $s_i$  after  $i$  steps when starting at the initial state  $s_0$  and following policy  $\pi$

The expected discounted reward of  $\pi$  is defined as:

$$Q(s, \pi) = \sum_{i=0}^{\infty} \gamma^i \left( \sum_{s_i} R(s_i) P(S_i = s_i \mid \pi, S_0 = s_0) \right).$$

and

$$U^\pi(s) = Q(s, \pi(s))$$

## Finding optimal strategies

The maximum expected utility of starting in state  $s$  is:

$$U^*(s) = \max_{\pi} Q(s, \pi) = \max_{\pi} \sum_{i=0}^{\infty} \gamma^i \left( \sum_{S_i} R(S_i) P(S_i \mid \pi, S_0 = s_0) \right).$$

In any state we choose the action maximizing the expected utility:

$$\delta(s) = \arg \max_a \sum_{s' \in \text{sp}(S)} P(s' \mid s, a) U^*(s')$$

with


$$U^*(s) = R(s) + \gamma \max_a \sum_{s' \in \text{sp}(S)} P(s' \mid s, a) U^*(s').$$

↪ we need to compute the value function.

But how do we calculate this?

# Value iteration

Start with an initial guess at the utility function, and iteratively refine this:



The updating function (a Bellman update):

$$U^{j+1}(s) := R(s) + \max_a \sum_{s'} P(s' | a, s) U^j(s').$$

# Value iteration: the algorithm

- ① Choose an  $\epsilon > 0$  to regulate the stopping criterion.
- ② Let  $U^0$  be an initial estimate of the utility function (for example, initialized to zero for all states).
- ③ Set  $i := 0$ .
- ④ **Repeat**
  - ① Let  $i := i + 1$ .
  - ② **For** each  $s \in \text{sp}(S)$

$$U^i(s) := R(s) + \gamma \cdot \max_a \sum_{s' \in \text{sp}(S)} P(s' | a, s) U^{i-1}(s').$$

- ⑤ **Until** Stopping criteria met (e.g.  $U^i(s) - U^{i-1}(s) < \epsilon$ , for all  $s \in \text{sp}(S)$ )

# Value iteration: an example

# Value iteration: the impact of the discounting factor

	1	2	3
1	7.91	8.9	10
2	6.82	6.79	2.62
3	5.83	5.66	4.85

The utility function for  $\gamma = 0.9$   
(51 iterations)

	1	2	3
1	→	→	
2	↑	↑	↑
3	↑	↑	←

The optimal strategy for  $\gamma = 0.9$

	1	2	3
1	-0.01	0.9	10
2	-0.1	-0.11	-4.29
3	-0.11	-0.11	-0.11

The utility function for  $\gamma = 0.1$  (17 iterations)

	1	2	3
1	→	→	
2	↑	↑	↑
3	↑	←	←

The optimal strategy for  $\gamma = 0.1$



## Value iteration: convergence

So the algorithm converges for this particular example, but does this hold in general?

**Yes**, it can be proven that there is only one “true” utility function and value iteration is guaranteed to converge to this utility function. Moreover,

$$m = \frac{\log(\epsilon(1 - \gamma)/2R_{\max})}{\log(\gamma)}$$

is an upper bound on the number of iterations required to achieve an error less than  $\epsilon$ .

## Beyond Value Iteration

There are other algorithms beyond value iteration:

- Policy Iteration: Instead of converging to the optimal value function  $V^*$  and then extracting the optimal policy, try to converge by changing the policy in each iteration.
- Heuristic Search: Value Iteration enumerates the entire state space before starting to solve the algorithm (like naive implementations of Dijkstra's algorithm). As in  $A^*$ , we can reduce the portion of the state-space that we consider by using admissible heuristics.
- Online decision making: There are also algorithms that, instead of computing the entire policy in advance, try to decide what action to perform in the current state.

# Summary

In general, in a **Markov decision process**:

- the world is *fully observable*, i.e., the agent can observe the true state of the world at any point in time,
- the uncertainty in the system is a result of the consequences of the actions being non-deterministic (when performing an action we make a state transition with a certain probability, but independent of the time step), and
- for each decision we get a reward (which may be negative) that may depend on the current world state but is independent of the time step.

**Value Iteration** can be used to find the optimal policy of any MDP.

# Reading

- *Chapter 9 Planning with Uncertainty* from the book "Artificial Intelligence: Foundations of Computational Agents" (2nd edition); in particular subchapters:
  - 9.1. Preferences and Utility
  - 9.2. One-Off Decisions
  - 9.3. Sequential Decisions
  - 9.4. The value of information and control
  - 9.5. Decision Processes

# Policy iteration

Instead of updating the utility function, make an initial guess at the optimal policy and perform an iterative refinement of this guess:

The **updating function**:

$$\pi_{i+1}(s) := \arg \max_a \sum_{s' \in \text{sp}(S)} P(s' | a, s) U_{\pi_i}(s').$$

The **evaluation function**:

$$U_{\pi_i}(s) = R(s) + \gamma \sum_{s' \in \text{sp}(S)} P(s' | \pi_i(s), s) U_{\pi_i}(s'),$$

which defines a system of linear equalities; the solution is  $U_{\pi_i}$ .

# Policy iteration: the algorithm

- ① Let  $\pi_0$  be an initial randomly chosen policy.
- ② Set  $i := 0$ .
- ③ **Repeat**
  - ① Find the utility function  $U_{\pi_i}$  corresponding to the policy  $\pi_i$  [Policy evaluation].
  - ② Let  $i := i + 1$ .
  - ③ **For** each  $s \in \text{sp}(S)$

$$\pi_i(s) := \arg \max_a \sum_{s' \in \text{sp}(S)} P(s' | a, s) U_{\pi_{i-1}}(s') [\text{Policy updating}].$$

- ④ **Until**  $\pi_i = \pi_{i-1}$

Policy evaluation can be performed both exactly ( $O(n^3)$ ) or approximately using a simplified version of value iteration.